Seminar II: Deep Learning-based Natural Language Processing

"BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding"
Paper review

Contents

1. Introduction
   - Contextualized word representations - ELMo
   - OpenAI GPT-1
   - Key Contribution
2. Methodology
   - BERT: Masked Language model
   - BERT: Next sentence prediction
   - BERT: Input representation
   - BERT: Model and training details
3. Results and Discussion
4. Conclusions

Introduction. Contextualized word representations - ELMo

Words have several meanings based on the context.

Contextualized word-embeddings can give words different embeddings based on the meaning they carry in the context of the sentence.
Source: https://jalammar.github.io/illustrated-bert/

Introduction. Contextualized word representations - ELMo

Bi-directional language modeling

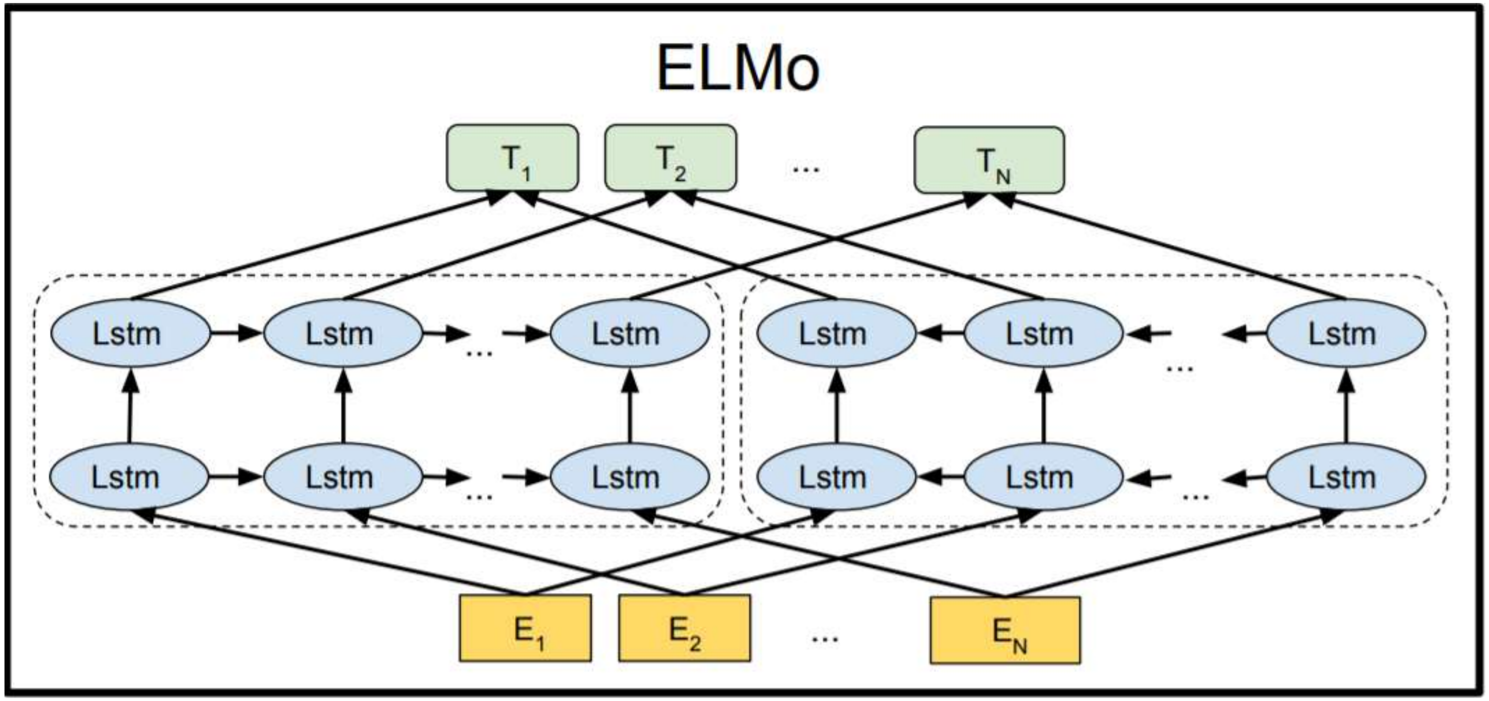

Bi-directional stacked LSTM.
Source: https://medium.com/saarthi-ai/elmo-for-contextual-word-embedding-for-text-classification-24c9693b0045

Introduction. OpenAI GPT-1

Transformer decoder without Enc-Dec attention

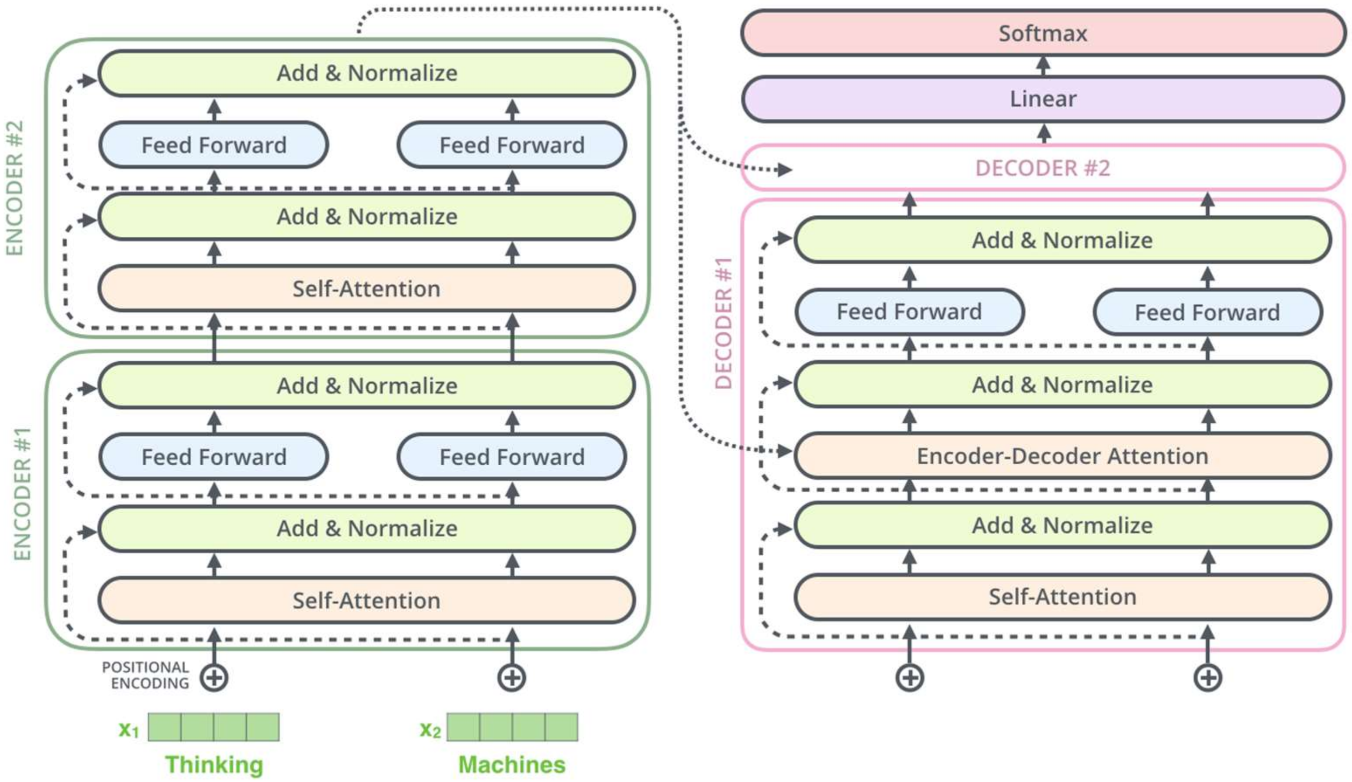

Source: https://jalammar.github.io/illustrated-transformer/

Introduction. OpenAI GPT-1

Pre-train on BooksCorpus

- 12 layers, 768 hidden size, 12 attention heads (110M parameters)

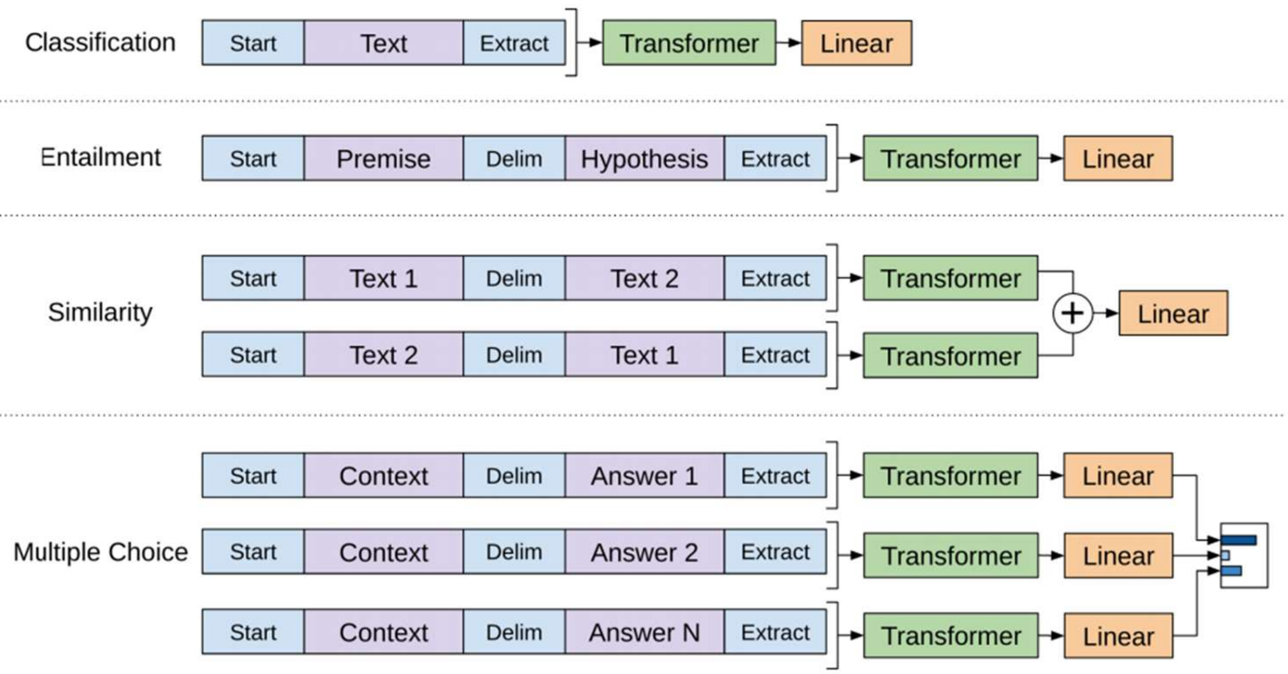

Source: Improving Language Understanding by Generative Pre-Training

Introduction. OpenAI GPT-1

Method	MNLI-m	MNLI-mm	SNLI	SciTail	QNLI	RTE
ESIM + ELMo [44] (5x)	-	-	89.3	-	-	-
CAFE [58] (5x)	80.2	79.0	89.3	-	-	-
Stochastic Answer Network [35] (3x)	80.6	80.1	-	-	-	-
CAFE [58]	78.7	77.9	88.5	83.3		
GenSen [64]	71.4	71.3	-	-	82.3	59.2
Multi-task BiLSTM + Attn [64]	72.2	72.1	-	-	82.1	**61.7**
Finetuned Transformer LM (ours)	**82.1**	**81.4**	**89.9**	**88.3**	**88.1**	56.0

Experimental results on natural language inference tasks
Source: Improving Language Understanding by Generative Pre-Training

Method	Story Cloze	RACE-m	RACE-h	RACE
val-LS-skip [55]	76.5	-	-	-
Hidden Coherence Model [7]	77.6	-	-	-
Dynamic Fusion Net [67] (9x)	-	55.6	49.4	51.2
BiAttention MRU [59] (9x)	-	60.2	50.3	53.3
Finetuned Transformer LM (ours)	**86.5**	**62.9**	**57.4**	**59.0**

Results on question answering and common-sense reasoning
Source: Improving Language Understanding by Generative Pre-Training

Introduction. Key Contribution

OBJECTIVE

To improve **the fine-tuning based LM approaches**

1. **Bidirectional approach – “Masked LM” (MLM) <- ELMo**
2. **Transformer – “GPT”**

Methodology – BERT: Masked Language model

Solution: Mask out k% of the input words, and then predict the masked words

- We always use k = 15%

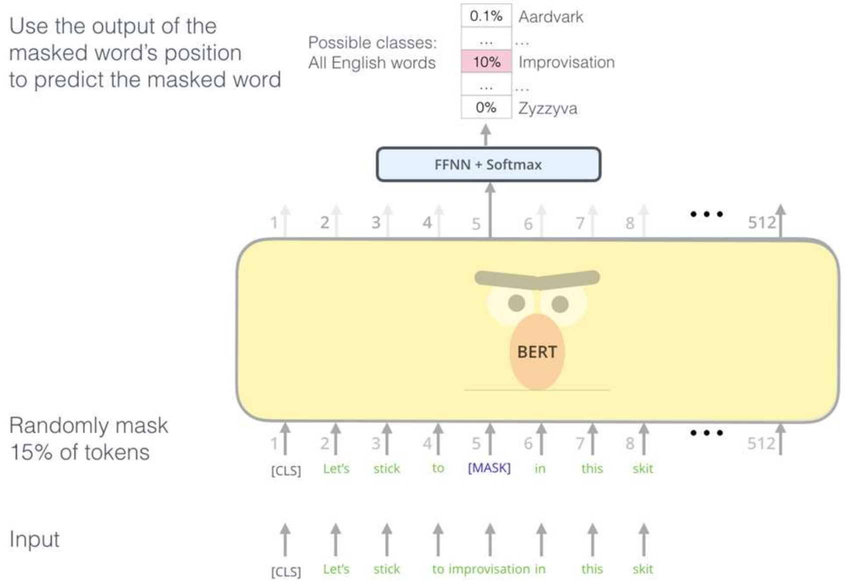

BERT's clever language modeling task masks 15% of words in the input and asks the model to predict the missing word
Source: https://jalammar.github.io/illustrated-bert/

Methodology – BERT: Masked Language model

Problem: Mask token never seen at fine-tuning

Solution: 15% of the words are *[mask]* but do not always replace with it. Instead:

- 80% of the time, replace with *[mask]*
  went to the store -> went to the *[mask]*
- 10% of the time, replace random word
  went to the store -> went to the *running*
- 10% of the time, keep same
  went to the store -> went to the *store*

Methodology – BERT: Next sentence prediction

Helpful for some downstream tasks

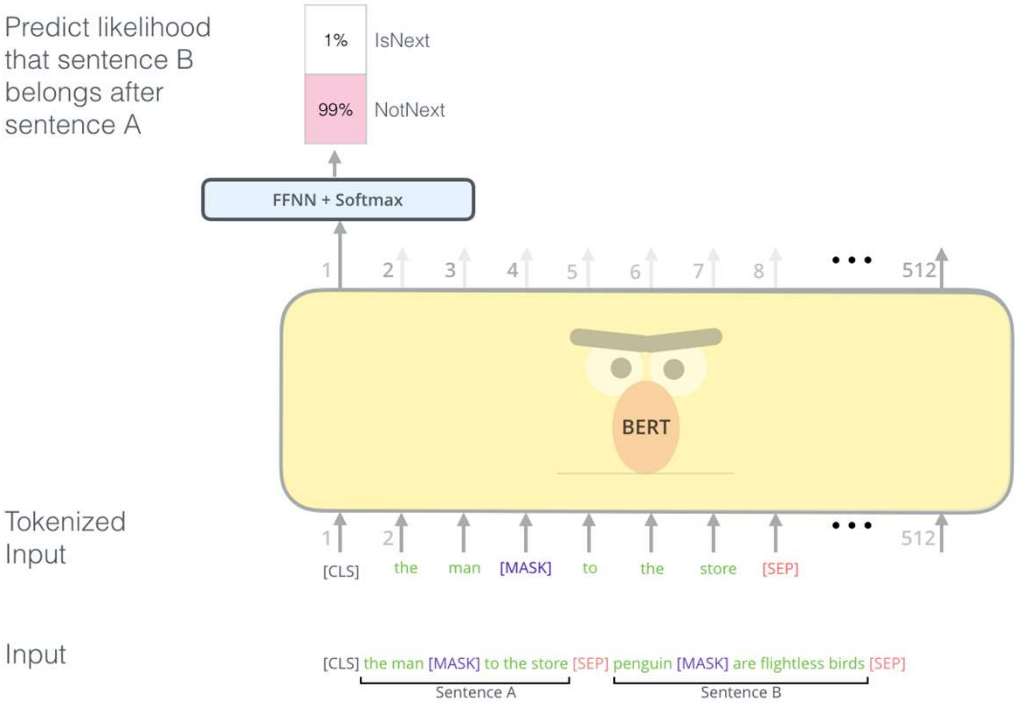

The second task BERT is pre-trained on is a two-sentence classification task. The tokenization is oversimplified in this graphic as BERT actually uses WordPieces as tokens rather than words --- so some words are broken down into smaller chunks.
Source: https://jalammar.github.io/illustrated-bert/

Methodology – BERT: Input representation

Use 30,000 WordPiece vocabulary on input

Each token is sum of three embeddings

Trainable Segment Embedding

CLS: Special tokens representing all input

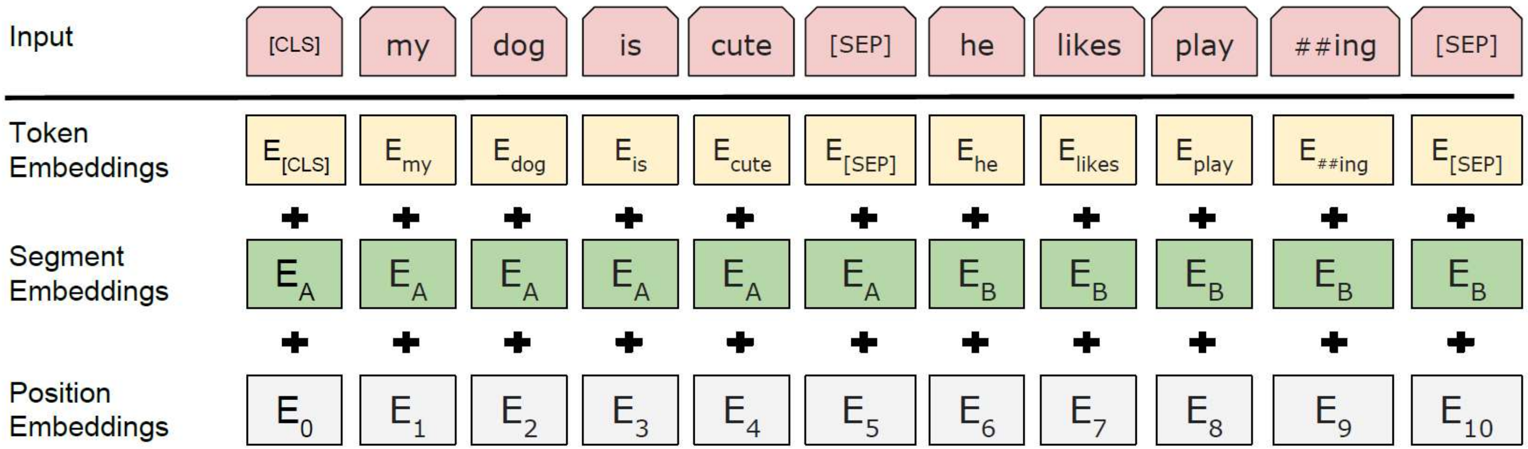

BERT input representation. The input embeddings are the sum of the token embeddings, the segmentation embeddings and the position embeddings.
Source: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Methodology – BERT: Model and training details

Pre-Train:

Data: Wikipedia (2.5B words) + BookCorpus (800M words)

BERT-base (Layer=12, Hidden=768, Head=12, Total Parameters= 110M)

BERT-large (Layer=24, Hidden=1024, Head=16, Total Parameters=340M)

Fine-Tune:

Different input/output

for different downstream task

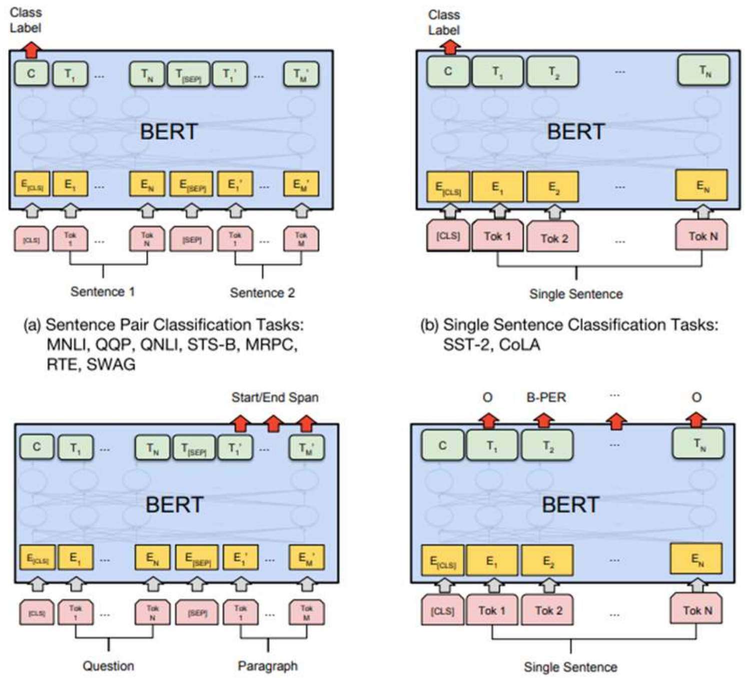

(a) Sentence Pair Classification Tasks: MNLI, QQP, QNLI, STS-B, MRPC, RTE, SWAG

(b) Single Sentence Classification Tasks: SST-2, CoLA

(c) Question Answering Tasks: SQuAD v1.1

(d) Single Sentence Tagging Tasks: CoNLL-2003 NER

Results and Discussions

System	MNLI-(m/mm)	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	**Average**
	392k	363k	108k	67k	8.5k	5.7k	3.5k	2.5k	-
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
$BERT_{BASE}$	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
$BERT_{LARGE}$	**86.7/85.9**	**72.1**	**92.7**	**94.9**	**60.5**	**86.5**	**89.3**	**70.1**	**82.1**

GLUE Test results (Collection of NLP tasks)
Source: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Rank	Name	Model	URL	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI	RTE	WNLI	AX
1	JDExplore d-team	Vega v1		91.3	73.8	97.9	94.5/92.6	93.5/93.1	76.7/91.1	92.1	91.9	96.7	92.4	97.9	51.4
2	Microsoft Alexander v-team	Turing NLR v5		91.2	72.6	97.6	93.8/91.7	93.7/93.3	76.4/91.1	92.6	92.4	97.9	94.1	95.9	57.0
3	DIRL Team	DeBERTa + CLEVER		91.1	74.7	97.6	93.3/91.1	93.4/93.1	76.5/91.0	92.1	91.8	96.7	93.2	96.6	53.3
4	ERNIE Team - Baidu	ERNIE		91.1	75.5	97.8	93.9/91.8	93.0/92.6	75.2/90.9	92.3	91.7	97.3	92.6	95.9	51.7
5	AliceMind & DIRL	StructBERT + CLEVER		91.0	75.3	97.7	93.9/91.9	93.5/93.1	75.6/90.8	91.7	91.5	97.4	92.5	95.2	49.1
6	DeBERTa Team - Microsoft	DeBERTa / TuringNLRv4		90.8	71.5	97.5	94.0/92.0	92.9/92.6	76.2/90.8	91.9	91.6	99.2	93.2	94.5	53.2
7	HFL iFLYTEK	MacALBERT + DKM		90.7	74.8	97.0	94.5/92.6	92.8/92.6	74.7/90.6	91.3	91.1	97.8	92.0	94.5	52.6
8	PING-AN Omni-Sinitic	ALBERT + DAAF + NAS		90.6	73.5	97.2	94.0/92.0	93.0/92.4	76.1/91.0	91.6	91.3	97.5	91.7	94.5	51.2
9	T5 Team - Google	T5		90.3	71.6	97.5	92.8/90.4	93.1/92.8	75.1/90.6	92.2	91.9	96.9	92.8	94.5	53.1
10	Microsoft D365 AI & MSR AI & GATECH	MT-DNN-SMART		89.9	69.5	97.5	93.7/91.6	92.9/92.5	73.9/90.2	91.0	90.8	99.2	89.7	94.5	50.2

GLUE Benchmark Leaderboard
Source: https://gluebenchmark.com/leaderboard

Results and Discussions

System	Dev EM	Dev F1	Test EM	Test F1
Top Leaderboard Systems (Dec 10th, 2018)				
Human	-	-	82.3	91.2
#1 Ensemble - nlnet	-	-	86.0	91.7
#2 Ensemble - QANet	-	-	84.5	90.5
Published				
BiDAF+ELMo (Single)	-	85.6	-	85.8
R.M. Reader (Ensemble)	81.2	87.9	82.3	88.5
Ours				
$\text{BERT}_{\text{BASE}}$ (Single)	80.8	88.5	-	-
$\text{BERT}_{\text{LARGE}}$ (Single)	84.1	90.9	-	-
$\text{BERT}_{\text{LARGE}}$ (Ensemble)	85.8	91.8	-	-
$\text{BERT}_{\text{LARGE}}$ (Sgl.+TriviaQA)	**84.2**	**91.1**	**85.1**	**91.8**
$\text{BERT}_{\text{LARGE}}$ (Ens.+TriviaQA)	**86.2**	**92.2**	**87.4**	**93.2**

SQuAD 1.1 results

Source: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Paragraph: The scientific revolution was a period when European ideas in classical Physics, Astronomy, Biology, Human Anatomy, Chemistry, and other classical sciences were rejected and led to doctrines supplanting those that had prevailed from ancient Greece to the middle ages which would lead to a transition to modern science. this period saw a fundamental transformation in scientific ideas across Physics, Astronomy, and Biology, in institutions supporting scientific investigation, and in the more widely held picture of the universe. individuals started to question all manners of things and it was this questioning that led to the scientific revolution, which in turn formed the foundations of contemporary sciences and the establishment of several modern scientific fields.

Question: What did the scientific revolution cause?

Answer: a transition to modern science

System	Dev EM	Dev F1	Test EM	Test F1
Top Leaderboard Systems (Dec 10th, 2018)				
Human	86.3	89.0	86.9	89.5
#1 Single - MIR-MRC (F-Net)	-	-	74.8	78.0
#2 Single - nlnet	-	-	74.2	77.1
Published				
unet (Ensemble)	-	-	71.4	74.9
SLQA+ (Single)	-		71.4	74.4
Ours				
$\text{BERT}_{\text{LARGE}}$ (Single)	78.7	81.9	80.0	83.1

SQuAD 2.0 results

Source: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Results and Discussions

System	Dev	Test
ESIM+GloVe	51.9	52.7
ESIM+ELMo	59.1	59.2
OpenAI GPT	-	78.0
$\text{BERT}_{\text{BASE}}$	81.6	-
$\text{BERT}_{\text{LARGE}}$	**86.6**	**86.3**
Human (expert)†	-	85.0
Human (5 annotations)†	-	88.0

SWAG Dev and Test accuracies

Tasks	Dev Set MNLI-m (Acc)	QNLI (Acc)	MRPC (Acc)	SST-2 (Acc)	SQuAD (F1)
$\text{BERT}_{\text{BASE}}$	84.4	88.4	86.7	92.7	88.5
No NSP	83.9	84.9	86.5	92.6	87.9
LTR & No NSP	82.1	84.3	77.5	92.1	77.8
+ BiLSTM	82.1	84.1	75.7	91.6	84.9

Ablation over the pre-training tasks

Hyperparams				Dev Set Accuracy		
#L	#H	#A	LM (ppl)	MNLI-m	MRPC	SST-2
3	768	12	5.84	77.9	79.8	88.4
6	768	3	5.24	80.6	82.2	90.7
6	768	12	4.68	81.9	84.8	91.3
12	768	12	3.99	84.4	86.7	92.9
12	1024	16	3.54	85.7	86.9	93.3
24	1024	16	3.23	86.6	87.8	93.7

Ablation over BERT model size

System	Dev F1	Test F1
ELMo (Peters et al., 2018a)	95.7	92.2
CVT (Clark et al., 2018)	-	92.6
CSE (Akbik et al., 2018)	-	**93.1**
Fine-tuning approach		
$\text{BERT}_{\text{LARGE}}$	96.6	92.8
$\text{BERT}_{\text{BASE}}$	96.4	92.4
Feature-based approach ($\text{BERT}_{\text{BASE}}$)		
Embeddings	91.0	-
Second-to-Last Hidden	95.6	-
Last Hidden	94.9	-
Weighted Sum Last Four Hidden	95.9	-
Concat Last Four Hidden	96.1	-
Weighted Sum All 12 Layers	95.5	-

CoNLL-2003 Named Entity Recognition results

Conclusions

- It works well.